

Yes. I think you should now answer these open questions decisively — otherwise the architecture risks drifting into endless calibration.

And importantly:
some of these are not merely engineering questions.

They are:

- ontology questions,
- governance questions,
- invariant questions.

Below is what I would answer.

Q1 — VAIDM provenance

Answer: computed in-house

I strongly recommend:

VAIDM must remain endogenous to PolygonEye/PolyAgora.

Not an external feed.

Why?

Because VAIDM is not merely:

- a macro indicator.

It is:

- a semantic deformation operator,
- a corridor-pressure estimator,
- a topology modifier.

If outsourced:

- the ontology fragments,
- governance no longer remains unified.

Correct architecture:

$$VAIDM_t = f(MacroPanel, CrossAssetStress, RealizedPnL, Fragmentation, TransitionDyna$$

YES:

- start from macro panel,
- evolve later toward richer topology.

Good current assumption.

Q2 — Universe

Answer: 13 partner futures

Strongly agree with current assumption.

Why?

Because:

- futures are liquid,
- corridor behavior observable,
- recoverability measurable,
- transition geometry clearer.

AGUR strategies are:

- already transformed objects,
- partially opaque,
- downstream abstractions.

PolyAgora should govern:

primary geometry first.

Not:

- already-packaged strategy outputs.

You can later map:

- strategies onto futures manifold.

But foundational Layer-1:

should remain:

- futures-native.

Correct choice.

Q3 — LEAST_BAD / BUFFER

This is VERY important

Their current draft is close but incomplete.

I would sharpen:

BUFFER

BUFFER is NOT merely:

- anti-whipsaw cooldown.

It is philosophically:

corridor ambiguity state.

Meaning:

- no stable local star yet,
- fragmentation unresolved,
- transition law still deforming,
- premature commitment dangerous.

Thus:

- freeze leverage,
- freeze aggressive convexity,
- preserve optionality.

BUFFER =

“do not collapse possibility space too early.”

This is deeply MRTP.

LEAST_BAD

LEAST_BAD is:

degraded admissibility without full rupture.

Meaning:

- no attractive corridor,
- but one corridor still survivable.

Not:

- opportunity state.

But:

- constrained continuity state.

Thus:

- minimal deployment,
- preserve continuity,
- avoid collapse,
- maintain re-entry.

This is very important because:

LEAST_BAD is not bearishness.

It is:

constrained maneuverability.

Huge distinction.

Q4 — Leverage cap

Answer: keep 1.25×

Strongly recommend:

DO NOT exceed this early.

Reason:

Your moat currently comes from:

- recoverability,
- low forced deleveraging risk.

Too much Kelly expansion:

- destroys ontology,
- destroys crisis edge,
- converts system into leveraged momentum allocator.

You need:

- asymmetry,
- not leverage.

Meaning:

- recoverable convexity,
- NOT:
- gross exposure escalation.

1.25× is enough initially.

Maybe later:

- dynamic regime-sensitive cap.

But not now.

Q5 — Regression tolerance band

My recommendation

This is extremely important.

Without this:

- Phase-II can silently destroy V7.10.

I would propose:

Metric	Allowed degradation
Sharpe	max -0.07

Metric	Allowed degradation
Max DD	max +1.5 pp
Corr SPY	max +0.03
Crisis windows	NO degradation allowed
Recovery latency	max +15%

MOST IMPORTANT:

Crisis moat is sacred.

Meaning:

If:

- GFC,
- COVID,
- 2022

materially degrade,
the architecture FAILS —
even if CAGR rises.

This must be explicit.

Q6 — θ_{ERQ} calibration

This is NOT merely empirical

This is one of the deepest points.

I would explicitly warn the CTO:

θ_{ERQ} is not a performance threshold.
It is a philosophical admissibility boundary.

Too low:

- system admits irreversible convexity,
- moat dies.

Too high:

- system reproduces V7.10 underparticipation.

My suggestion:

DO NOT begin from:

- return optimization.

Begin from:

- recoverability falsification.

Meaning:

- identify structures that historically destroyed maneuverability,
- ensure ERQ rejects them FIRST.

Then gradually widen admissibility.

This is critical.

Q7 — MRTP coefficients

This is the deepest unresolved issue

Because these coefficients define:

the ethical geometry of the system.

Especially:

$$\eta$$

which governs:

- preference for recoverable convexity.

My recommendation:

Initially bias strongly toward:

- recoverability,
- corridor width,
- fragmentation avoidance.

Meaning:

$$\gamma, \delta > \eta$$

early.

Then gradually increase:

$$\eta$$

only after:

- crisis moat proven intact.

This is crucial.

Otherwise:

- the system will prematurely become convexity-seeking.

The deepest recommendation overall

Your CTO team must NEVER forget:

PolyAgora's edge is NOT superior prediction.

It is:

superior admissibility governance under changing market ontology.

That distinction must remain sacred.

Because the temptation in Phase-II will be:

- more convexity,
- more beta,
- more participation,
- more Kelly,
- more return chasing.

That temptation can quietly destroy:

- MRTP,
- recoverability,
- and the Fixed-Star logic itself.

Final guidance to CTO

I would summarize to them:

Core Governance Constraint

Phase-II may improve convex participation,
but it must NEVER invert the hierarchy:

$$\textit{Recoverability} > \textit{Convexity} > \textit{Return}$$

If at any point:

- convexity dominates recoverability,
or:
- return dominates admissibility,

the architecture ceases to be PolyAgora
and collapses into conventional hedge-fund optimization.

The crisis moat is not a side effect.

It is the manifestation of the invariant itself.

Phase-II succeeds ONLY if:

- recoverability remains primary,
 - admissible convexity improves,
 - and future maneuverability survives across regime transformation.
-

That is the correct answer to these open questions.