Yes — **V7.7 is a real improvement in governance, but not yet in alpha extraction.**

What the dashboard shows:

V7.6α base is still the strongest full-return line:
Sharpe 0.82 | CAGR 3.23% | Max DD -10.25%

V7.7 full stack gives:
Sharpe 0.65 | CAGR 1.89% | Max DD -7.65%

So V7.7 **reduces drawdown and volatility**, but it also cuts too much upside.

The best new component is actually **V7.7-ADD only**:
Sharpe 0.80 | CAGR 2.74% | Max DD -8.61%

That is important. It means:

> **ADD-lite is useful. Kelly is too conservative in current form.**

My evaluation:

1. ADD-lite works.
It deforms risk exposure without destroying the signal. This is exactly what an external polygon should do: qualify, compress, warn, but not replace the engine.

2. Kelly overlay is over-defensive.
The full V7.7 stack moves too much into CASH. It improves survival, but kills compounding. Current Kelly is acting more like a brake than a capital-growth optimizer.

3. V7.6α remains the benchmark to beat.
V7.7 should not be presented as "better returns." It should be presented as:
same architecture + better crisis survivability, but Kelly needs recalibration.

4. The dashboard itself is good.
Clear, CTO-friendly, investor-readable. Strong sections: equity curve, β timeline, drawdown, weights, and explanatory tooltips.

My CTO instruction:

Keep **V7.7-ADD** as the serious candidate.
Treat **V7.7-Kelly/full stack** as experimental.

Next test:

> V7.6α + ADD-lite + softer Kelly
> Kelly only activates when ADD/recoverability field confirms stable convex opportunity, not continuously.

In short:

Breakthrough: ADD-lite validated.
Problem: Kelly too blunt.
Next version: conditional Kelly, not permanent Kelly.